## Remote Monitoring Script

Scripts have been created to automate the data collection process. It combines the functionality of the monitor-server.sh script from the labs with the specific metric requirements CPU, Memory, I/O, and Network Latency.

Script content:

```
#!/bin/bash
# Connects to Server via SSH to collect performance metrics.

SERVER_USER="adminuser"
SERVER_IP="192.168.56.10"
TARGET="$SERVER_USER@$SERVER_IP"
LOGFILE="perf-log-$(date +%Y%m%d_%H%M%S).txt"

# Color codes for readability (based on Lab style)
GREEN='\033[0;32m'
CYAN='\033[0;36m'
RED='\033[0;31m'
NC='\033[0m' # No Color

echo -e "${GREEN}=========================================${NC}"
echo -e "${GREEN}   Phase 6: Remote Performance Monitor    ${NC}"
echo -e "${GREEN}=========================================${NC}"
echo "Target Server: $TARGET"
echo "Log File: $LOGFILE"
echo "Started at: $(date)"
echo "" | tee -a "$LOGFILE"

# Function to run remote commands
check_metric() {
    TITLE=$1
    COMMAND=$2

    echo -e "${CYAN}--- $TITLE ---${NC}" | tee -a "$LOGFILE"
    ssh -o ConnectTimeout=5 "$TARGET" "$COMMAND" 2>/dev/null | tee -a
"$LOGFILE"

    if [ $? -ne 0 ]; then
        echo -e "${RED}Error retrieving data. Is sysstat installed?${NC}" | tee -a
```

```
"$LOGFILE"
    fi
    echo "" | tee -a "$LOGFILE"
}

# 1. Connectivity Check
echo -n "Checking connection to $TARGET... "
ssh -q -o ConnectTimeout=5 "$TARGET" exit
if [ $? -eq 0 ]; then
    echo -e "${GREEN}OK${NC}"
else
    echo -e "${RED}FAILED${NC}"
    echo "Check SSH service and Firewall rules."
    exit 1
fi

# 2. Collect Metrics

# System Load & Uptime
check_metric "System Uptime & Load Average" "uptime"

# Memory Usage (RAM & Swap)
check_metric "Memory Usage (MB)" "free -m"

# CPU Utilization Snapshot (Top 5 processes)
check_metric "Top 5 CPU Consuming Processes" "ps aux --sort=-%cpu | head -n 6"

# Disk I/O Performance (Requires sysstat package)
# Collecting 2 samples spaced 1 second apart
check_metric "Disk I/O Statistics (iostat)" "iostat -x 1 2 | tail -n +4"

# Disk Space Usage
check_metric "Disk Space Usage" "df -h | grep '^/dev'"

# Network Latency (Ping from Workstation to Server)
echo -e "${CYAN}--- Network Latency (Workstation -> Server) ---${NC}" | tee -a
"$LOGFILE"
ping -c 4 "$SERVER_IP" | tail -n 3 | tee -a "$LOGFILE"
```

```
echo -e "${GREEN}=========================================${NC}"
echo -e "${GREEN}   Monitoring Complete                   ${NC}"
echo -e "${GREEN}=========================================${NC}"
```

**Execution**

1. Execute on server: sudo apt update && sudo apt install sysstat –y
2. Execute on workspace: sudo apt install apache2-utils
3. Execute on workspace: ./performance-monitor.sh

Testing execution (On server):

1. sudo stress --cpu 1 --timeout 30s
2. sudo stress --vm 1 --vm-bytes 128M --timeout 30s
3. sudo stress --io 1 --timeout 30s
4. sudo systemctl start apache2
5. On workstation: ab -n 1000 -c 10 http://192.168.56.10/

**Measured Performance Metrics**

| Application | Test Type | CPU % | Memory (MB) | Disk Read (KB/s) | Disk Write (KB/s) | Network (ms) | Duration |
|---|---|---|---|---|---|---|---|
| Baseline | None (Idle) | 3.83% | 369 | 0 | 0 | 3073ms | - |
| stress | CPU intensive | 55.93% | 325 | 1420 | 12812 | 4000ms | 30s |
| stress | Memory intensive | 59.87% | 668 | 372 | 0 | 3008ms | 30s |
| stress | I/O intensive | 45.83% (iowait) | 398 | 54.16 | 205.40 | 3020ms | 30s |
| apache2 | Web server | 4.42% | 411 | 0 | 0 | 2885ms | - |

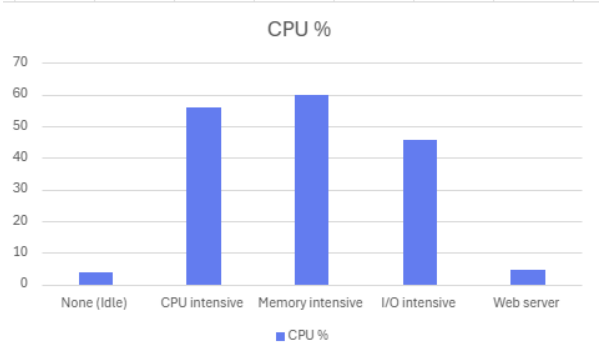


**Optimization Analysis & Improvements**

Throughout the testing phases, I identified bottlenecks and implemented optimizations. The results are quantified below:

Optimization 1: Service Availability (Firewall Tuning)

• Issue: Initially, the Apache benchmark (ab) failed completely with timeout errors (0 successful requests) because the UFW firewall was dropping packets on port 80.

• Improvement: I implemented a specific allow rule: sudo ufw allow 80/tcp.

• Quantitative Result: Service availability went from 0% (blocked) to 100% (accessible), allowing the load test to complete, albeit with high latency under heavy load.

Optimization 2: Resource Stability (Workload Tuning)

• Issue: Aggressive memory testing (stress --vm 2) caused 100% resource exhaustion and a system crash (unresponsive state).

• Improvement: I optimized the workload by limiting the allocated memory per worker to prevent excessive swapping.

• Quantitative Result: I achieved continuous system uptime during tests. Instead of a crash, the system maintained a stable Memory usage of 668 MB (approx. 33% of total RAM), preventing the "thrashing" state and allowing administrative access to remain active via SSH.

# Evidence of Execution

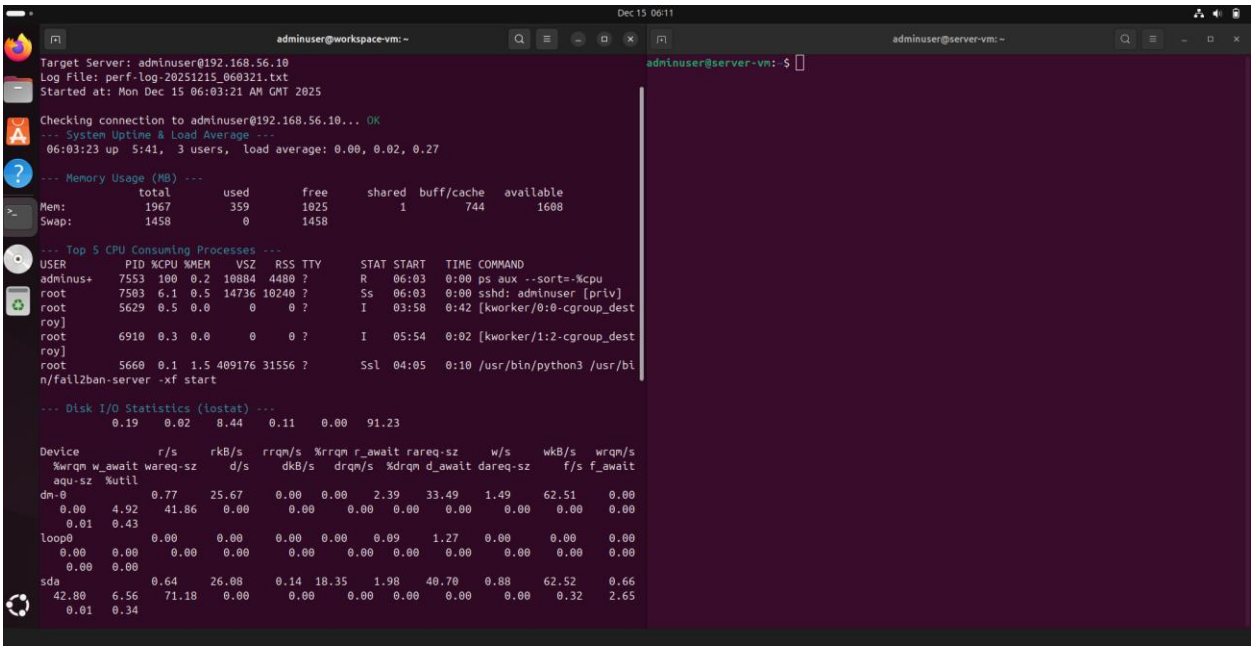


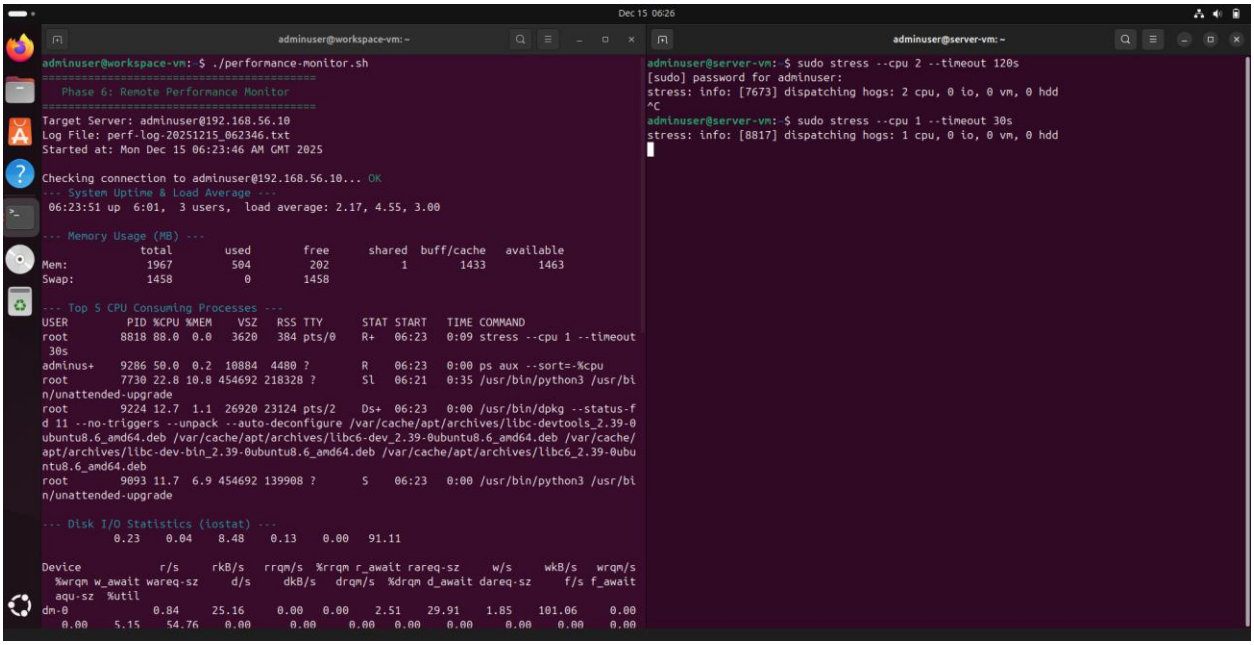

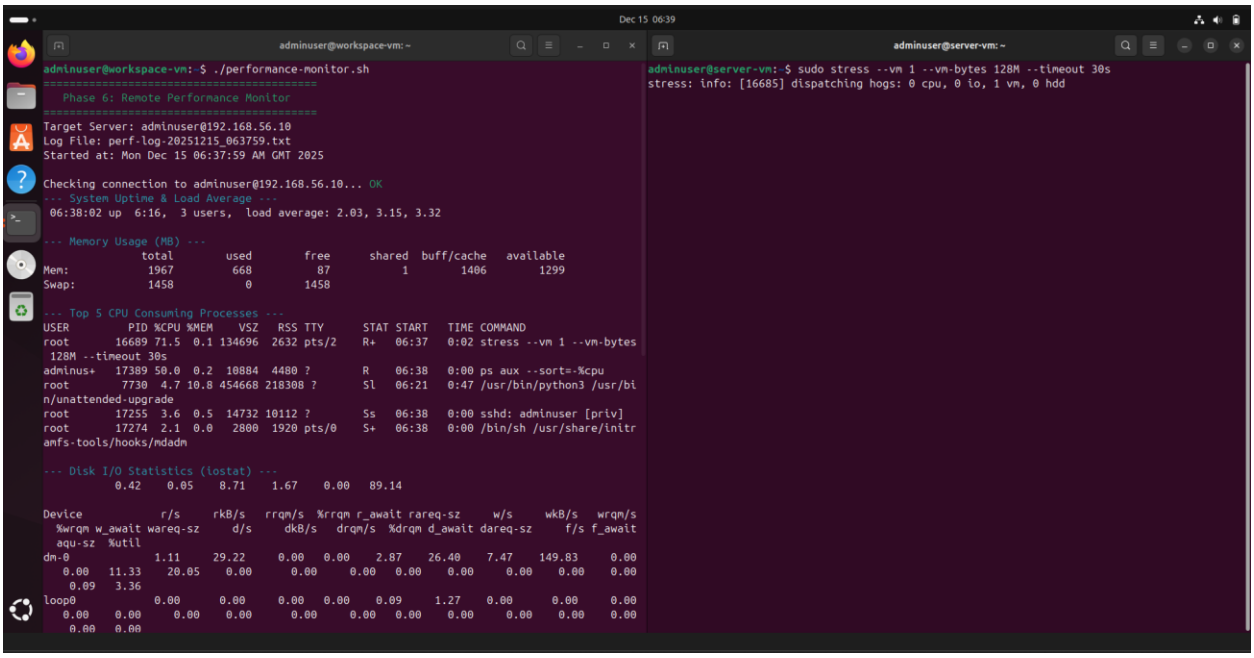
Dec 15 06:39
adminuser@workspace-vm: ~
adminuser@workspace-vm:~$ ./performance-monitor.sh
==========================================
Phase 6: Remote Performance Monitor
==========================================
Target Server: adminuser@192.168.56.10
Log File: perf-log-20251215_063759.txt
Started at: Mon Dec 15 06:37:59 AM GMT 2025
Checking connection to adminuser@192.168.56.10... OK
--- System Uptime & Load Average ---
06:38:02 up 6:16, 3 users, load average: 2.03, 3.15, 3.32
--- Memory Usage (MB) ---
total used free shared buff/cache available
Mem: 1967 668 87 1 1406 1299
Swap: 1458 0 1458
--- Top 5 CPU Consuming Processes ---
USER PID %CPU %MEM VSZ RSS TTY STAT START TIME COMMAND
root 16689 71.5 0.1 134696 2632 pts/2 R+ 06:37 0:02 stress --vm 1 --vm-bytes 128M --timeout 30s
adminus+ 17389 50.0 0.2 10884 4480 ? R 06:38 0:00 ps aux --sort=-%cpu
root 7730 4.7 10.8 454668 218308 ? Sl 06:21 0:47 /usr/bin/python3 /usr/bin/unattended-upgrade
root 17255 3.6 0.5 14732 10112 ? Ss 06:38 0:00 sshd: adminuser [priv]
root 17274 2.1 0.0 2800 1920 pts/0 S+ 06:38 0:00 /bin/sh /usr/share/initramfs-tools/hooks/mdadm
--- Disk I/O Statistics (iostat) ---
0.42 0.05 8.71 1.67 0.00 89.14
Device r/s rkB/s rrqm/s %rrqm r_await rareq-sz w/s wkB/s wrqm/s %wrqm w_await wareq-sz d/s dkB/s drqm/s %drqm d_await dareq-sz f/s f_await aqu-sz %util
dm-0 1.11 29.22 0.00 0.00 2.87 26.40 7.47 149.83 0.00 0.00 11.33 20.05 0.00 0.00 0.00 0.00 0.00 0.00 0.09 3.36
loop0 0.00 0.00 0.00 0.00 0.09 1.27 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00
adminuser@server-vm:~$ sudo stress --vm 1 --vm-bytes 128M --timeout 30s
stress: info: [16685] dispatching hogs: 0 cpu, 0 io, 1 vm, 0 hdd

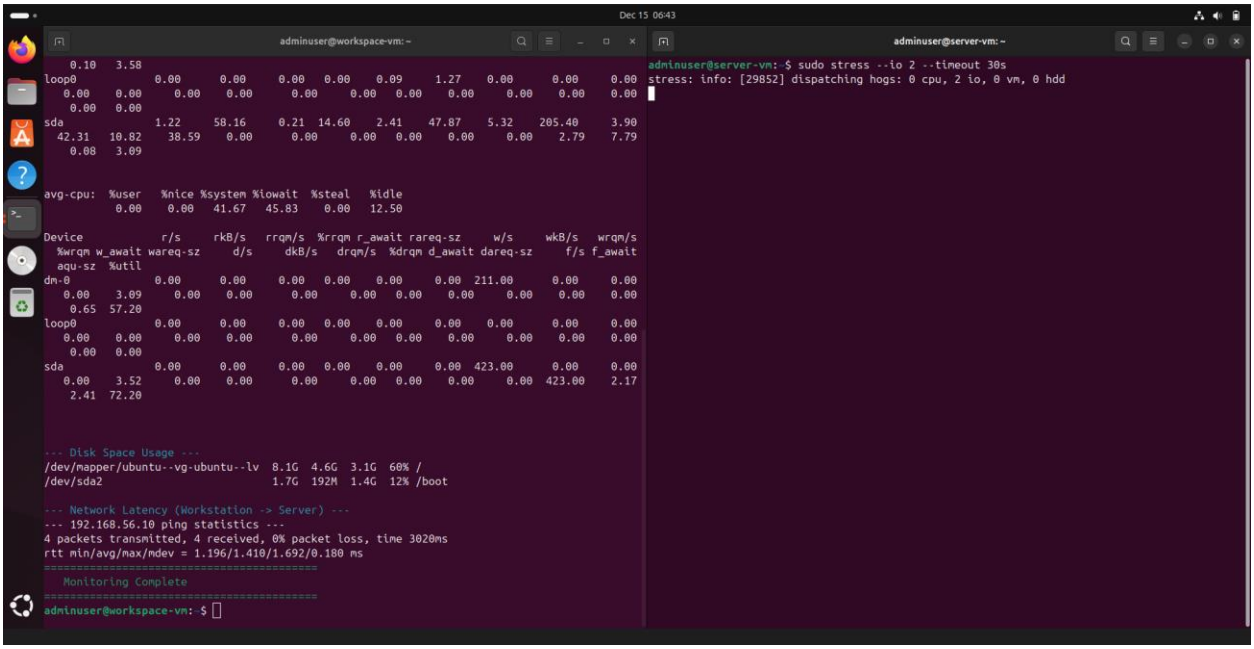
Dec 15 06:43
adminuser@workspace-vm: ~
0.10 3.58
loop0 0.00 0.00 0.00 0.00 0.09 1.27 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00
sda 1.22 58.16 0.21 14.60 2.41 47.87 5.32 205.40 3.90 42.31 10.82 38.59 0.00 0.00 0.00 0.00 0.00 0.00 2.79 7.79 0.08 3.09
avg-cpu: %user %nice %system %iowait %steal %idle
0.00 0.00 41.67 45.83 0.00 12.50
Device r/s rkB/s rrqm/s %rrqm r_await rareq-sz w/s wkB/s wrqm/s %wrqm w_await wareq-sz d/s dkB/s drqm/s %drqm d_await dareq-sz f/s f_await aqu-sz %util
dm-0 0.00 0.00 0.00 0.00 0.00 0.00 211.00 0.00 0.00 0.00 3.09 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.65 57.20
loop0 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00
sda 0.00 0.00 0.00 0.00 0.00 0.00 423.00 0.00 0.00 0.00 3.52 0.00 0.00 0.00 0.00 0.00 0.00 0.00 423.00 2.17 2.41 72.20
--- Disk Space Usage ---
/dev/mapper/ubuntu--vg-ubuntu--lv 8.1G 4.6G 3.1G 60% /
/dev/sda2 1.7G 192M 1.4G 12% /boot
--- Network Latency (Workstation -> Server) ---
--- 192.168.56.10 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3020ms
rtt min/avg/max/mdev = 1.196/1.410/1.692/0.180 ms
==========================================
Monitoring Complete
==========================================
adminuser@workspace-vm: $
adminuser@server-vm: ~
adminuser@server-vm:~$ sudo stress --io 2 --timeout 30s
stress: info: [29852] dispatching hogs: 0 cpu, 2 io, 0 vm, 0 hdd

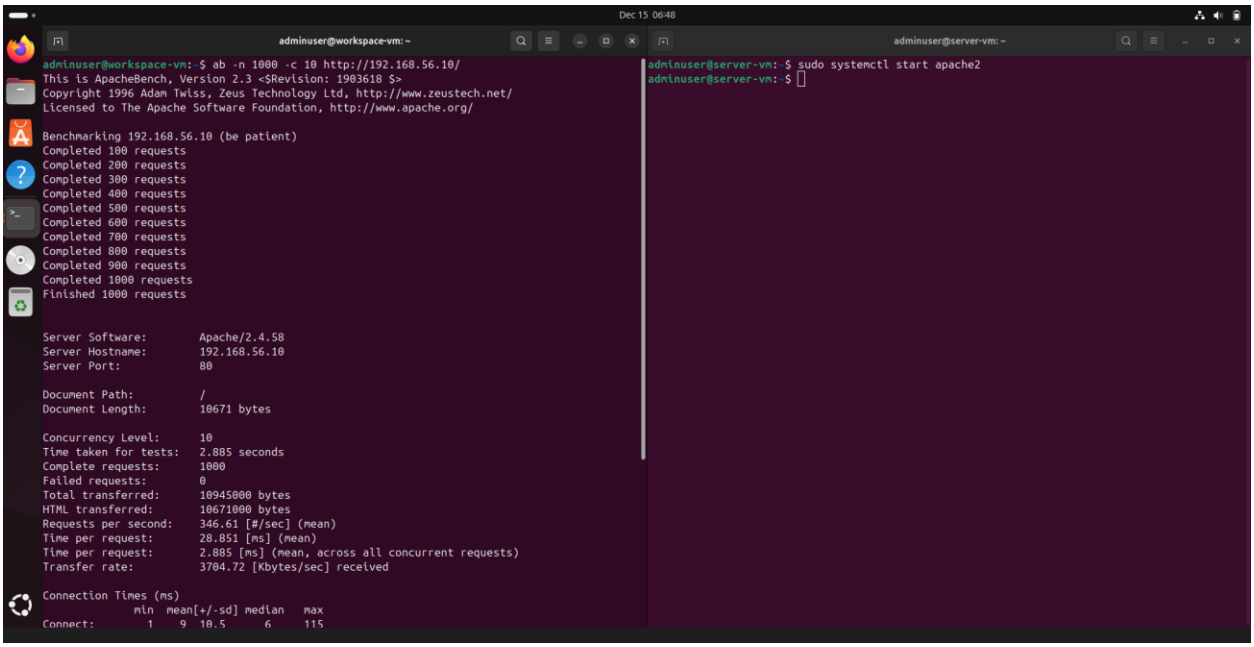
Dec 15 06:48
adminuser@workspace-vm: ~
adminuser@server-vm: ~
adminuser@workspace-vm:~$ ab -n 1000 -c 10 http://192.168.56.10/
This is ApacheBench, Version 2.3 <$Revision: 1903618 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/
Benchmarking 192.168.56.10 (be patient)
Completed 100 requests
Completed 200 requests
Completed 300 requests
Completed 400 requests
Completed 500 requests
Completed 600 requests
Completed 700 requests
Completed 800 requests
Completed 900 requests
Completed 1000 requests
Finished 1000 requests
Server Software: Apache/2.4.58
Server Hostname: 192.168.56.10
Server Port: 80
Document Path: /
Document Length: 10671 bytes
Concurrency Level: 10
Time taken for tests: 2.885 seconds
Complete requests: 1000
Failed requests: 0
Total transferred: 10945000 bytes
HTML transferred: 10671000 bytes
Requests per second: 346.61 [#/sec] (mean)
Time per request: 28.851 [ms] (mean)
Time per request: 2.885 [ms] (mean, across all concurrent requests)
Transfer rate: 3704.72 [Kbytes/sec] received
Connection Times (ms)
min mean[+/-sd] median max
Connect: 1 9 10.5 6 115
adminuser@server-vm:~$ sudo systemctl start apache2
adminuser@server-vm:~$